

Applied Data Analysis (CS401)



Lecture 11
Handling text data
Part II
26 Nov 2025

EPFL

Maria Brbic

Announcements

- To all ADAmericans:
Happy Thanksgiving!
- Homework H2 due this **today** EOD
 - Reminder: We won't answer questions asked during final 24h
- Projects:
 - Milestone P2 grades have been released
 - **Project Milestone P3 released today!**
- Friday's lab session:
 - Exercises on handling text ([Exercise 10](#))



Feedback

Give us feedback on this lecture here:

<https://go.epfl.ch/ada2025-lec11-feedback>

- What did you (not) like about this lecture?
- What was (not) well explained?
- On what would you like more (fewer) details?
- ...

Let me open my **bag of tricks for bags of words** for you! But only if you were good children...

Recap



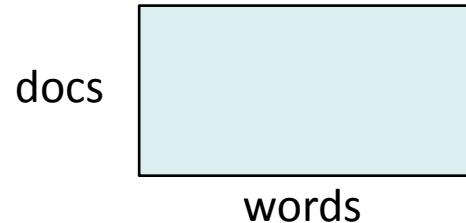
Reminder:
bag-of-words matrix

docs

words

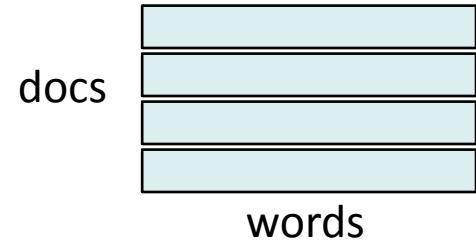
Revisiting the 4 typical tasks

- Document retrieval
 - Document classification
 - Sentiment analysis
 - Topic detection
-
- TF-IDF matrix to the rescue
 - Entry for doc d , word w :
$$\text{tf}(d, w) * \text{idf}(w)$$



Typical task 1: document retrieval

- Nearest-neighbor method in spirit of kNN
- Compare query doc q to all documents in the collection (i.e., rows of the TF-IDF matrix)
- Rank docs from collection in increasing order of distance
- Distance metrics
 - Typically cosine distance ($= 1 - \text{cosine similarity}$)
 - Recall: cosine similarity of q and $v = \langle q / \|q\|, v / \|v\| \rangle$
 - If rows are L2-normalized, may simply take dot products $\langle q, v \rangle$

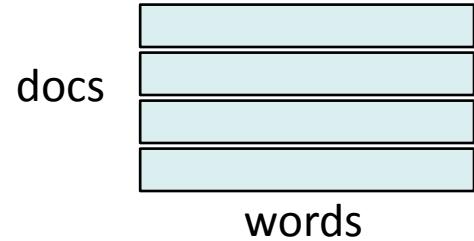


Typical task 1: document retrieval

- This is just the most basic approach
- For efficiency
 - Start by filtering documents by presence of query terms (use efficient full-text index)
 - Hugely narrows down set of documents to be ranked
- Google et al. do much more...
 - Query-independent relevance: PageRank
 - Boost recent results
 - Personalization, contextualization
 - ...

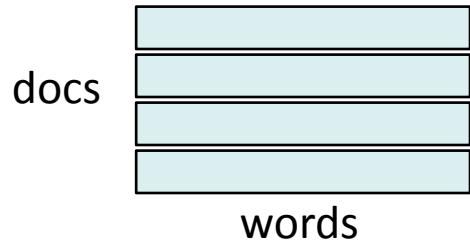
Typical task 2: document classification

- Use TF-IDF matrix as feature matrix for supervised methods (cf. lecture 7)
- Often more features (words) than documents
- What's the danger with this?
- High model capacity can lead to overfitting (high variance)
- Potential solutions:
 - Use more data (i.e., more labeled training docs)
 - Decrease model capacity:
 - Feature selection
 - Regularization (two slides from now)
 - Dimensionality reduction (a few slides from now)
 - Use ensemble methods such as random forests



Typical task 3: sentiment analysis

- When treated as classification:
Ctrl-C Ctrl-V previous slide
- When treated as regression:
Pretty much the same (most supervised methods work for both classification and regression)

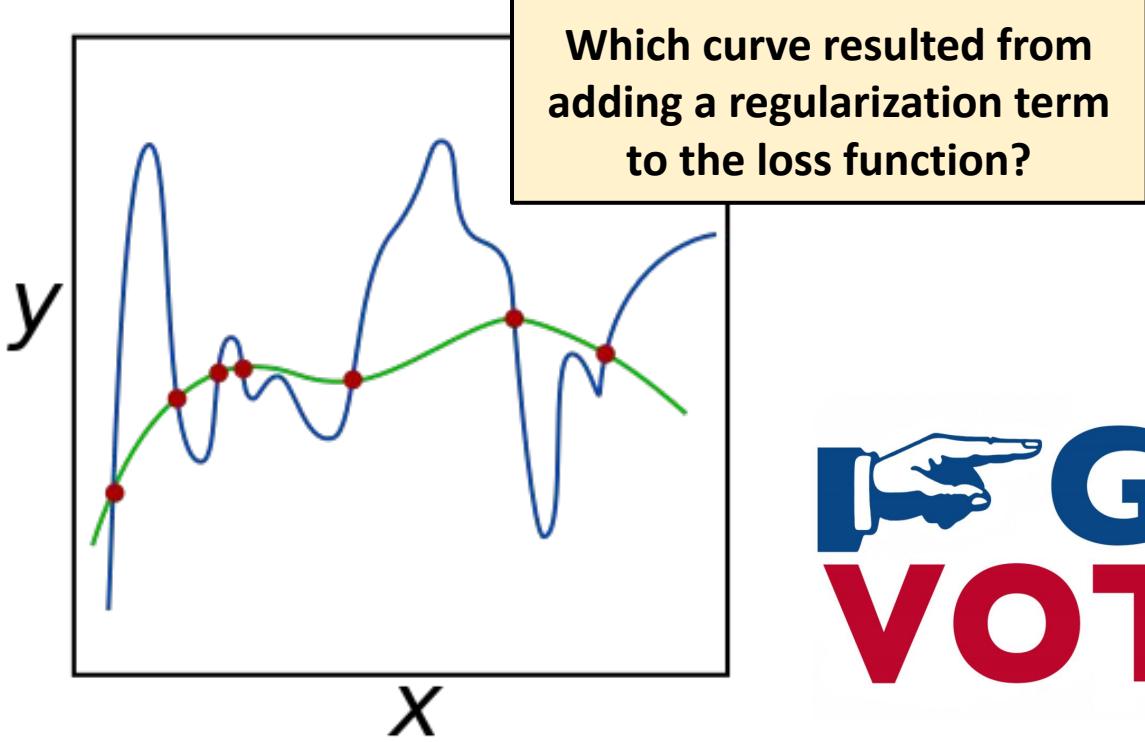


Regularization

- E.g., linear regression:
Find weight vector β that minimizes $\sum_{i=1}^n (y_i - \mathbf{x}_i^\top \beta)^2$
(\mathbf{x}_i : feature vector of i -th data point; y_i : label of i -th data point, e.g., sentiment [1–5 stars] expressed in document i)
- If one word j appears only in docs with sentiment 5, we can obtain very small training error on these docs by making β_j large enough
- But doesn't generalize to unseen test data!
- Remedy: penalize very large positive and very large negative weights:

$$\text{minimize } \sum_{i=1}^n (y_i - \mathbf{x}_i^\top \beta)^2 + \lambda \sum_{j=1}^p \beta_j^2$$

Regularization



POLLING TIME

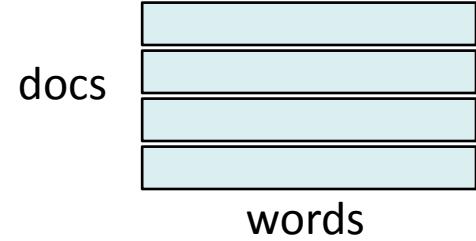
- Scan QR code or go to <https://go.epfl.ch/ada2025-lec11-poll>

 GO
VOTE



Typical task 4: topic detection

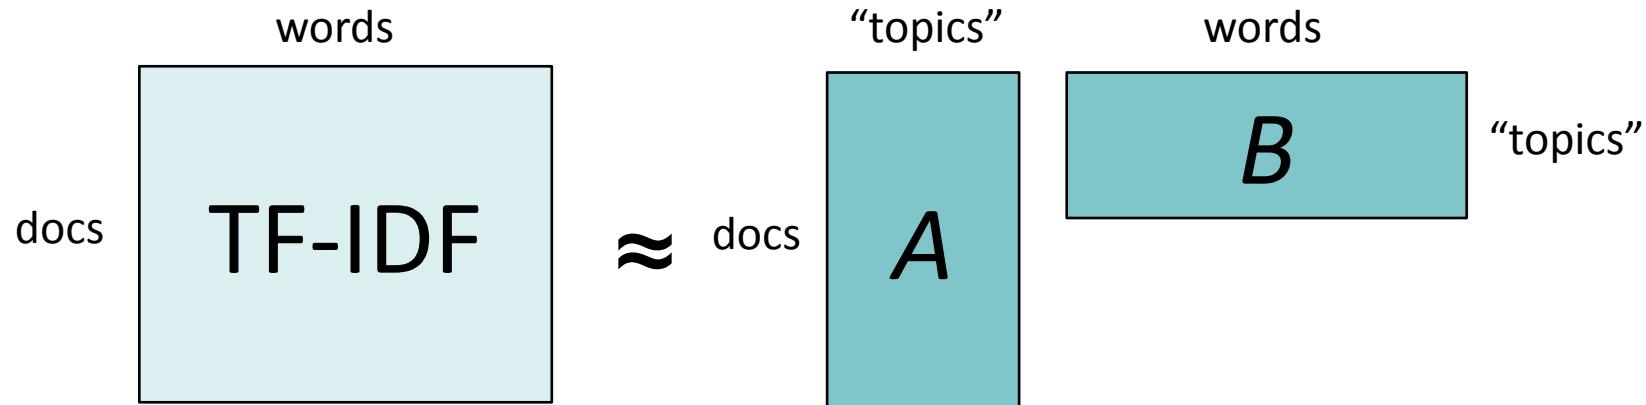
- Cluster rows of TF-IDF matrix (each row a data point)
- Manually inspect clusters and label them with descriptive names (e.g., “news”, “sports”, “romance”, “tech”, “politics”)
- In principle, may use k-means, k-medoids, etc.
- But can be difficult if dimensionality is large ($\#\text{words} \gg \#\text{docs}$)
 - “Curse of dimensionality”
 - Many outliers



What else can we do?

Typical task 4: topic detection

- Alternative approach: **matrix factorization**

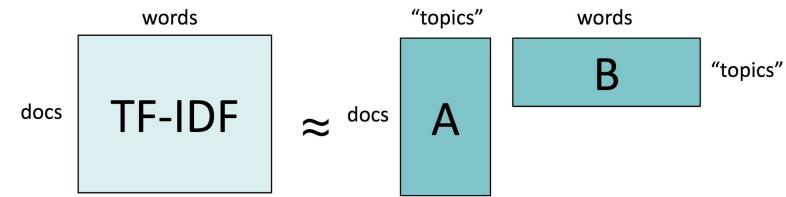


- Assume docs and words have representation in (latent) “topic space”
- (IDF-weighted) word frequency modeled as dot product of doc’s vectors and word’s vectors in topic space
- #topics \ll #words (\rightarrow “dimensionality reduction”): $D \cdot W \rightarrow (D+W) \cdot T$
- Topics interpretable in doc space (A’s cols) and word space (B’s rows)

Typical task 4: topic detection

- Optimization problem:

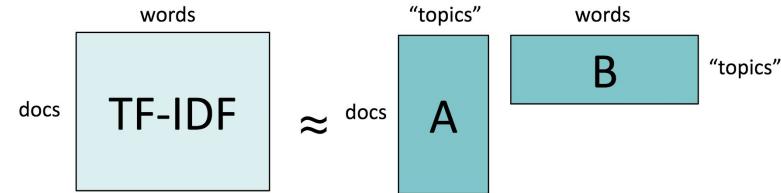
- Find A, B such that AB is as close to TF-IDF matrix as possible
- That is, minimize $\sum_{d=1}^N \sum_{w=1}^M (T_{dw} - A_d \cdot B_w)^2$



- This is called **latent semantic analysis (LSA)**

Typical task 4: topic detection

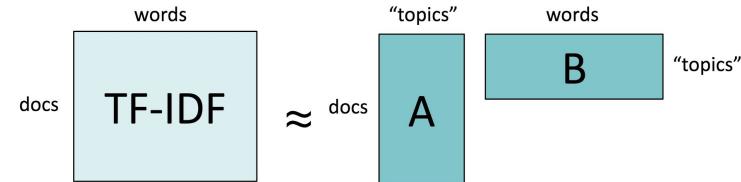
You already know how to efficiently compute this, from your linear algebra class: **singular-value decomposition** (SVD)



- $T = USV^T$
- Freebie: columns of U and V are orthonormal bases (yay!)
- S is diagonal (with values in decreasing order) and captures “importance” of topic (amount of variation in corpus w.r.t. topic)
- If you want k topics, keep only the first k columns of U and V , and the first k rows and columns of S
→ U', S', V'
- E.g., $A = U', B = S'V'^T$ or $A = U'S', B = V'^T$

Typical task 4: topic detection

- Recall potential problem with clustering and classification and regression: “curse of dimensionality”
- Matrix factorization via LSA solves these problems for you:
 - Use A instead of original TF-IDF matrix
 - That is, cluster (or learn to classify or regress) in topic space, rather than word space
- Topic representation from LSA is simply a vector, not a probability distribution over topics
- Probabilistic: LDA = Latent Dirichlet Allocation (p.t.o.)



Commercial break



ça roule pour moi

LDA: probabilistic topic modeling

- Latent Dirichlet Allocation (*not* Latent Discriminant Analysis!)
- Document := bag of words
- Topic := probability distribution over words
- Each document has a (latent) distribution over topics
- “Generative story” for generating a doc of length n :
 d := sample a topic distribution for the doc (\leftarrow “Dirichlet”)
for $i = 1, \dots, n$
 - t := sample a topic from topic distribution d
 - w := sample a word from topic t
 - Add w to the bag of words of the doc to be generated

Topics

gene	0.04
dna	0.02
genetic	0.01
...	

life	0.02
evolve	0.01
organism	0.01
...	

brain	0.04
neuron	0.02
nerve	0.01
...	

data	0.02
number	0.02
computer	0.01
...	

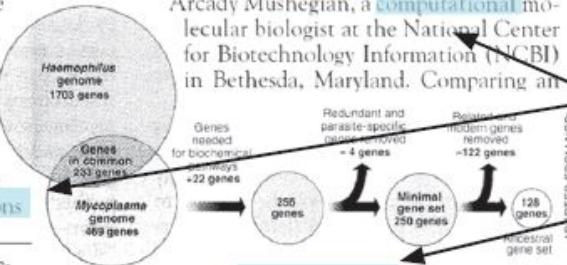
Documents

Seeking Life's Bare (Genetic) Necessities

COLD SPRING HARBOR, NEW YORK—How many genes does an organism need to survive? Last week at the genome meeting here,* two genome researchers with radically different approaches presented complementary views of the basic genes needed for life. One research team, using computer analyses to compare known genomes, concluded that today's organisms can be sustained with just 250 genes, and that the earliest life forms required a mere 128 genes. The other researcher mapped genes in a simple parasite and estimated that for this organism, 800 genes are plenty to do the job—but that anything short of 100 wouldn't be enough.

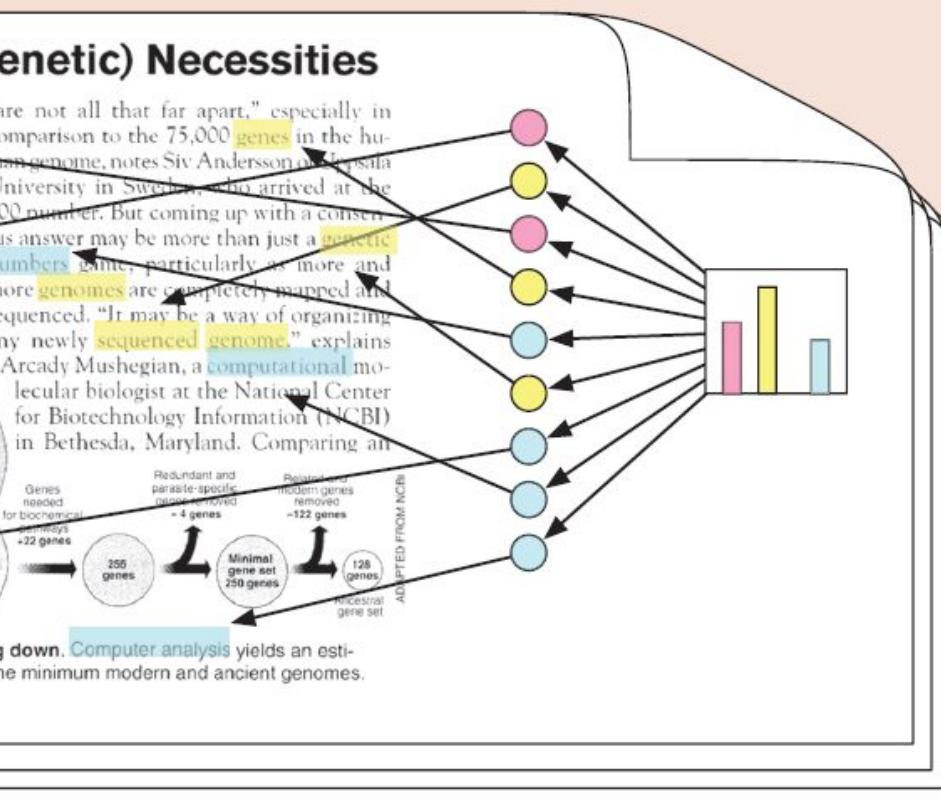
Although the numbers don't match precisely, those predictions

"are not all that far apart," especially in comparison to the 75,000 genes in the human genome, notes Siv Andersson of Uppsala University in Sweden, who arrived at the 800 number. But coming up with a consensus answer may be more than just a genetic numbers game, particularly as more and more genomes are completely mapped and sequenced. "It may be a way of organizing any newly sequenced genome," explains Arcady Mushegian, a computational molecular biologist at the National Center for Biotechnology Information (NCBI) in Bethesda, Maryland. Comparing an



Stripping down. Computer analysis yields an estimate of the minimum modern and ancient genomes.

Topic proportions and assignments

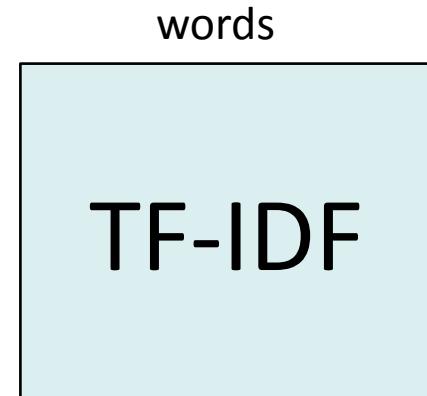


Topic inference in LDA

- LDA is unsupervised (topics come out “magically”)
- Input:
 - Docs represented as bags of words
 - Number K of topics
- Output:
 - K topics (distributions over words)
 - For each doc: distribution over K topics
- How is this done?
 - Find distributions (i.e., topics, docs) that maximize the likelihood of the observed documents (maximum likelihood)

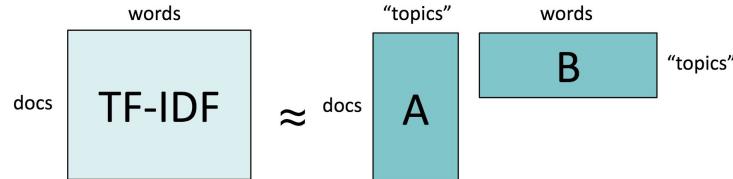
Question:

- “Which of these word pairs is more closely related?”
 - (car, bus)
 - (car, astronaut)
- How to quantify this?
- Detour:
 - How to quantify closeness of two docs?
 - E.g., cosine of **rows** of TF-IDF matrix
- Retour:
 - How to quantify closeness of two words?
 - E.g., cosine of **cols** of TF-IDF matrix



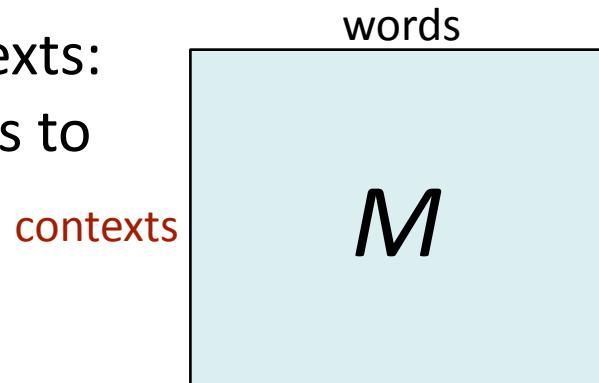
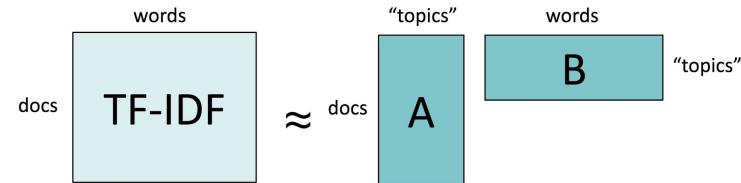
Sparsity in TF-IDF matrix

- Two docs (i.e., rows of TF-IDF matrix)
 - “Do you love men?”
 - “Adorest thou the likes of Adam?”
 - Cosine of row vectors of TF-IDF matrix == 0
- Same problem when comparing two words (i.e., cols of matrix)
- Solution:
 - Move from sparse to dense vectors
 - But how?
 - Latent semantic analysis (LSA)!
 - Use columns of B as dense vectors representing words



“Word vectors”

- Columns of TF-IDF matrix (sparse) or of word-by-topic matrix B (dense)
- Problem:
 - Entire doc treated as one bag of words
 - All information about word proximity, syntax, etc., is lost
- Solution:
 - Instead of full docs, consider local contexts: windows of L (e.g., 3) consecutive words to **left and right of the target word**
 - Rows of matrix: not docs, but contexts



“Word vectors”

contexts

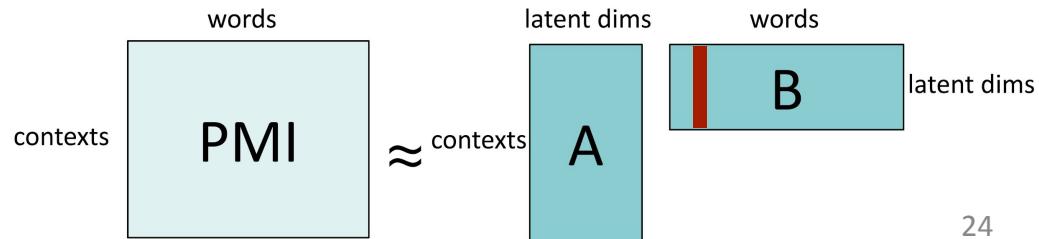
M

- What to use as entries of word/context matrix?
- Straightforward: same as TF-IDF, but with contexts as “pseudo-docs”: $M[c,w] = \text{TF-IDF}(c,w)$
- May use any other measures of statistical association
- E.g., pointwise mutual information (PMI):

$$M[c,w] = \text{PMI}(c,w) = \log \frac{\Pr(c, w)}{\Pr(c) \Pr(w)}$$

“How much more likely are c and w to occur together than if they were independent?”

- word2vec: factor PMI matrix and use columns of B as word vectors



Beyond bags of words

From words to texts

- Word vectors represent, well, words
- How to represent larger units, such as sentences, paragraphs, docs?
- Typical approach: take sum/average of word vectors
- Note: this is roughly also what bags of words are (when using “one-hot” encoding for words, i.e., vector with exactly one 1, rest 0)
- More recently: **learn** vectors for longer units
 - [Cr5](#), [sent2vec](#)
 - Convolutional neural networks
 - Recurrent neural networks, e.g., LSTM, [ELMo](#)
 - Transformer-based models, e.g., [BERT](#) (next slides), GPT-*

Contextualized word vectors

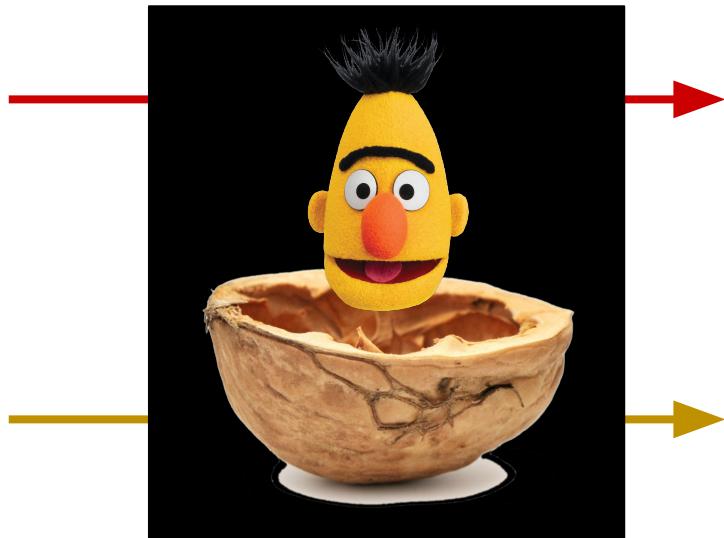
- Motivating example:
 - “The **bat** flew into the cave.”
 - vs. “He swung the **bat** and hit a home run.”
- Classic word vectors (e.g., word2vec) cannot distinguish these two cases; same vector used for both instances of “**bat**”
- Solution: contextualized word vectors
 - E.g., BERT



BERT in a nutshell

- Introduced in 2018 by Google Research

<START>
the
bat
flew
...



[1.00,0.70,0.90,0.50,0.06,...]
[0.54,0.75,0.56,0.45,0.09,...]
[0.44,0.76,0.77,0.31,0.82,...]
[0.91,0.62,0.53,0.75,0.74,...]
[0.92,0.37,0.25,0.49,0.24,...]

} doc vector
} context-
ualized
word
vectors

<START>
he
swung
the
bat
...

[0.85,0.62,0.71,0.11,0.58,...]
[0.49,0.25,0.22,0.36,0.75,...]
[0.61,0.87,0.73,0.96,0.52,...]
[0.58,0.02,0.01,0.92,0.76,...]
[0.53,0.42,0.64,0.26,0.01,...]

NLP pipeline

- Tokenization
- Sentence splitting
- Part-of-speech (POS) tagging
- Named-entity recognition (NER)
- Coreference resolution
- Parsing
 - Shallow parsing
(a.k.a. chunking)
 - Constituency parsing
 - Dependency parsing

Screenshot of the Stanford CoreNLP web interface showing the results of the NLP pipeline for the sentence: "Chase Manhattan and its merger partner J.P. Morgan and Citibank, which was involved in moving about \$100 million for Raul Salinas de Gortari, brother of a former Mexican president, to banks in Switzerland, are also expected to sign on."

Stanford CoreNLP

Output format: Visualise

Please enter your text here:

Chase Manhattan and its merger partner J.P. Morgan and Citibank, which was involved in moving about \$100 million for Raul Salinas de Gortari, brother of a former Mexican president, to banks in Switzerland, are also expected to sign on.

Part-of-Speech:

1 Chase Manhattan and its merger partner J.P. Morgan and Citibank, which was involved in moving about \$100 million for Raul Salinas de Gortari, brother of a former Mexican president, to banks in Switzerland, are also expected to sign on.

DOLLAR CD CD IN NNP NNP IN NNP IN CC NNP IN WDT VBD VBN IN VBG RB \$ 100 million for Raul Salinas de Gortari, brother of a former Mexican president, to banks in NNP VBP RB VBN TO VB IN Switzerland, are also expected to sign on.

Named Entity Recognition:

1 Chase Manhattan and its merger partner J.P. Morgan and Citibank, which was involved in moving about \$100 million for Raul Salinas de Gortari, brother of a former Mexican president, to banks in Switzerland, are also expected to sign on.

Organization Organization Org MONEY Person Location Location

Coreference:

1

Chase Manhattan and its merger partner J.P. Morgan and Citibank, which was involved in moving about \$100 million for Raul are also expected to sign on.

Mention Mention Mention Mention

Basic dependencies:

1 Chase Manhattan and its merger partner J.P. Morgan and Citibank, which was involved in moving about \$100 million for Raul are also expected to sign on.

dep poss cc conj cc conj cc nssubjpass nssubjpass auxpass prep pcomp nssubjpass prep pobj np

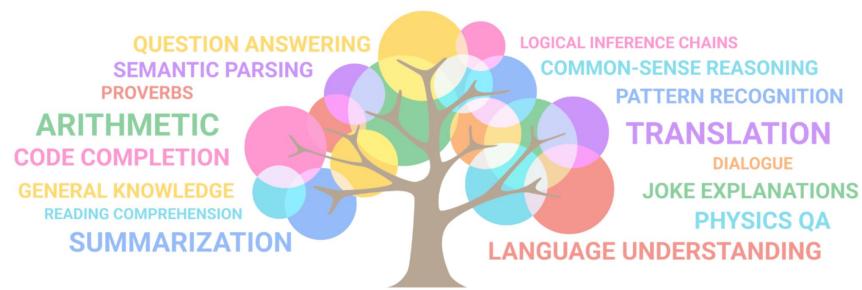
NLP pipeline

- Implemented by [Stanford CoreNLP](#), [nltk](#), [spaCy](#), etc.
- Sequential model
 - Fixed order of steps
 - Early errors will propagate downstream
 - Fixed order not optimal for all cases (e.g., syntax usually done before semantics, but semantics might be useful for inferring syntax)
- Hence, current research: learn all tasks jointly ([early example](#))
- To learn how all this magic is implemented
 - Take [CS-431](#) (Intro to NLP), [CS-552](#) (Modern NLP)

Today's trend: generative language models

The screenshot shows the OpenAI Playground interface at beta.openai.com/playground. The main area displays a text generation task: "Turn the following negative beer review into a positive review: '440ml can pour, 4.5% abv version. Aroma and taste of light grains, slight fruity hop undertone, clear golden appearance. Light bodied. Very average American style light lager.'". Below this, the generated positive review is shown: "440ml can pour, 4.5% abv version. Aroma of light grains and a slight fruity hop undertone. Clear golden appearance. Light bodied. A crisp and classic American style light lager. Refreshing and perfect for any occasion!". The interface includes standard browser controls, a user profile icon for 'epfl-dlab', and a sidebar with buttons for 'Save', 'View code', 'Share', and settings.

- E.g., OpenAI GPT-*, ChatGPT
- Input: text
- Output: text
- Many NLP tasks can be formulated in this framework, by “prompting” the language model with the right input



Feedback

Give us feedback on this lecture here:

<https://go.epfl.ch/ada2025-lec11-feedback>

- What did you (not) like about this lecture?
- What was (not) well explained?
- On what would you like more (fewer) details?
- ...